

Graph neural networks

Marek Dědič

TZN-05

Contents



- Motivation
- Tasks
- GNN
- DeepWalk
- node2vec
- GCN
- GraphSAGE

Motivation

- Motivation
- Tasks
- GNN
- DeepWalk
- node2vec
- GCN
- GraphSAGE

Motivation

- Graphs are a natural way to represent many different mathematical problems
- Examples: social networks, molecules, transportation networks, knowledge graphs, computer networks...



1

¹Veličković et al., 2020.


Definitions

- Graph $G = (V, E)$
- $ne[v]$ is the set of nodes connected to v by an edge
- $co[v]$ is the set of edges leading to or from v
- x_v are features of a node v
- x_e are features of an edge e

Tasks


- Motivation
- **Tasks**
- GNN
- DeepWalk
- node2vec
- GCN
- GraphSAGE

Node classification




2

Link prediction




³Kubara, 2020.

Graph learning



⁴Kubara, 2020.

Community detection



5

⁵Girvan et al., 2002.

- Motivation
- Tasks
- **GNN**
- DeepWalk
- node2vec
- GCN
- GraphSAGE

- We're solving “node classification”
- We adapt neural networks to graph data
- Formally, classification on the space $\mathcal{G} \times \mathcal{N}$ – pairs graph + one of its nodes
- We assume features $\mathbf{x}_v$ and $\mathbf{x}_e$ for each node v and edge e
- We model for each node v a hidden state $\mathbf{h}_v$ capturing information about v and its neighbourhood

- We model for each node v a hidden state $\mathbf{h}_v$ capturing information about v and its neighbourhood as


$$\mathbf{h}_v = f(\mathbf{x}_v, \mathbf{x}_{co[v]}, \mathbf{h}_{ne[v]}, \mathbf{x}_{ne[v]})$$

- ...and an output for node v as

$$\mathbf{o}_v = g(\mathbf{h}_v, \mathbf{x}_v)$$


- We can calculate the hidden state $\mathbf{h}_v$ iteratively for all nodes as

$$\mathbf{H}^{t+1} = F(\mathbf{H}^t, \mathbf{X})$$



- We can view the iterative application of function F as an RNN
- We may replace the iterative process with an MLP
- A GNN doesn't use edge hidden states, only node hidden states



GNN – modifications



- Motivation
- Tasks
- GNN
- **DeepWalk**
- node2vec
- GCN
- GraphSAGE

- Inspired by word2vec and skip-gram methods from NLP
- M random walks of length L from a given node to generate sequences of nodes, analogous to sentences


DeepWalk 2



8


⁸Perozzi et al., 2014.

DeepWalk 3



⁹Cohen, 2018.

DeepWalk – skip-gram



Skip-gram

10

¹⁰Mikolov et al., 2013.

- For each vector v we learn its representation $f : \mathcal{V} \rightarrow \mathbb{R}^{|\mathcal{V}| \times d}$
- We try to maximize the value

$$P(\{v_{i-w}, \dots, v_{i+w}\} \setminus v_i | f(v_i))$$

- Assuming independence, we simplify

$$P(\{v_{i-w}, \dots, v_{i+w}\} \setminus v_i | f(v_i)) = \prod_{\substack{j=i-w \\ j \neq i}}^{i+w} P(v_j | f(v_i))$$

- This is a problem solved by word2vec
- It's further necessary to optimize using hierarchical softmax

DeepWalk – conclusion

- Using ideas from word2vec to learn on graphs
- DeepWalk cannot classify new nodes – it is necessary to retrain the model

- Motivation
- Tasks
- GNN
- DeepWalk
- **node2vec**
- GCN
- GraphSAGE

- Also inspired by word2vec, but in a different way
- We are looking for an embedding $\text{node2vec} : \mathcal{G} \rightarrow \mathbb{R}^n$

- For each vector v we learn its representation $f : \mathcal{V} \rightarrow \mathbb{R}^{|\mathcal{V}| \times d}$
- We try to maximize the value

$$P(ne[v]|f(v))$$

- Assuming independence, we can simplify


$$P(ne[v]|f(v)) = \prod_{v_i \in ne[v]} P(v_i|f(v))$$

- It's further necessary to optimize using hierarchical softmax and negative sampling

- How to find $ne[v]$?
- 2 extrema: BFS, DFS
- Regular random walk: If I'm at node v , I go to some node connected to v with equal probability for all such nodes
- We modify the random walk to interpolate between BFS and DFS
- If coming to v from t , we assign weights to candidates for the next node x as

$$\alpha(t, x) = \begin{cases} \frac{1}{p} & \text{for } d_{tx} = 0 \\ 1 & \text{for } d_{tx} = 1 \\ \frac{1}{q} & \text{for } d_{tx} = 2 \end{cases}$$

where d_{tx} is the length of the shortest path from t to x and p, q are model parameters



11

¹¹Grover et al., 2016.

- Using ideas from word2vec to learn on graphs
- An improvement over DeepWalk
- node2vec cannot classify new nodes – it is necessary to retrain the model

node2vec – demo

- Motivation
- Tasks
- GNN
- DeepWalk
- node2vec
- **GCN**
- GraphSAGE

- We are solving “node classification”
- Inspired by CNNs

- Expressed as a neural network operating on the whole graph, similar to GNN
- Each layer has the form

$$\mathbf{H}^{l+1} = f\left(\mathbf{H}^{(l)}, \mathbf{A}\right)$$

where $\mathbf{A}$ is the adjacency matrix

- We can express a simple layer as


$$f\left(\mathbf{H}^{(l)}, \mathbf{A}\right) = \sigma\left(\mathbf{A}\mathbf{H}^{(l)}\mathbf{W}^{(l)}\right)$$

where $\mathbf{W}$ are the weights and σ is the activation function

- If we multiplied by $\mathbf{A}$, we would forget the information in the given node at each step. We solve this by using instead the matrix $\hat{\mathbf{A}} = \mathbf{A} + \mathbf{I}$
- Because the matrix $\mathbf{A}$ is not normalized, the scale would change after each layer. To prevent this, we use instead the symmetrically normalized matrix $\mathbf{D}^{-\frac{1}{2}}\mathbf{A}\mathbf{D}^{-\frac{1}{2}}$ where $\mathbf{D}$ is the diagonal matrix of degrees of individual nodes
- We thus obtain one layer of the network as

$$f\left(\mathbf{H}^{(l)}, \mathbf{A}\right) = \sigma\left(\hat{\mathbf{D}}^{-\frac{1}{2}}\hat{\mathbf{A}}\hat{\mathbf{D}}^{-\frac{1}{2}}\mathbf{H}^{(l)}\mathbf{W}^{(l)}\right)$$

GCN 4



12

¹²Kipf, 2016.

- Using ideas from CNNs to learn on graphs
- GCNs are a first-order approximation of localized spectral filters on graphs – see Kipf; Welling, 2017

- Motivation
- Tasks
- GNN
- DeepWalk
- node2vec
- GCN
- **GraphSAGE**

GraphSAGE 1

- GraphSAGE is an inductive framework for generating node embeddings – there is no need to have the whole graph during training
- We learn a representation for each node v using the representations of its neighbors, i.e.


$$\mathbf{h}_v^{(l+1)} = \sigma \left(\mathbf{W}^{(l)} \cdot \text{concat} \left(\mathbf{h}_v^{(l)}, \mathbf{h}_{ne[v]}^{(l)} \right) \right)$$

where

$$\mathbf{h}_{ne[v]}^{(l)} = \text{AGGREGATE}_l \left(\left\{ \mathbf{h}_u^{(l)} \mid u \in ne[v] \right\} \right)$$

- We need a special loss function

GraphSAGE 2



13

¹³Hamilton et al., 2017.

- We can choose an aggregator for each layer independently
- The mean aggregator yields a GCN-like layer
- An LSTM aggregator can capture more complex relationships (we need to randomly permute the neighbors)
- Pooling:

$$\text{AGGREGATE}_l = \max \left\{ \sigma \left(\mathbf{W}_{\text{pool}} \mathbf{h}_u^{(l)} + \mathbf{b} \right) \mid u \in \text{ne}[v] \right\}$$

GraphSAGE – conclusion

- An inductive approach is suitable for graphs where nodes are added after training
- We don't need to re-train the model, it will work with new data



Bibliography I

- COHEN, Elior, 2018. *node2vec: Embeddings for Graph Data* [Medium] [online]. 2018-04-23. [visited on 2020-11-03]. Available from:
<https://towardsdatascience.com/node2vec-embeddings-for-graph-data-32a866340fef>.
- GIRVAN, M.; NEWMAN, M. E. J., 2002. Community structure in social and biological networks. *Proceedings of the National Academy of Sciences* [online]. Vol. 99, no. 12, pp. 7821–7826 [visited on 2020-11-03]. ISSN 0027-8424, ISSN 1091-6490. Available from DOI: 10.1073/pnas.122653799. Publisher: National Academy of Sciences Section: Physical Sciences.
- GORI, M.; MONFARDINI, G.; SCARSELLI, F., 2005. A new model for learning in graph domains. In: *Proceedings. 2005 IEEE International Joint Conference on Neural Networks, 2005*. Vol. 2, 729–734 vol. 2. Available from DOI: 10.1109/IJCNN.2005.1555942. ISSN: 2161-4407.
- GROVER, Aditya; LESKOVEC, Jure, 2016. node2vec: Scalable feature learning for networks. In: *Proceedings of the 22nd ACM SIGKDD international conference on Knowledge discovery and data mining*, pp. 855–864.
- HAMILTON, Will; YING, Zhitao; LESKOVEC, Jure, 2017. Inductive representation learning on large graphs. In: *Advances in neural information processing systems*, pp. 1024–1034.

Bibliography II

- KIPF, Thomas N., 2016. *How powerful are Graph Convolutional Networks?* [online]. 2016-09-30. [visited on 2020-11-04]. Available from: <http://tkipf.github.io/graph-convolutional-networks/>.
- KIPF, Thomas N.; WELLING, Max, 2017. Semi-Supervised Classification with Graph Convolutional Networks. In: *5th International Conference on Learning Representations, {ICLR} 2017, Toulon, France, April 24-26, 2017, Conference Track Proceedings*. Toulon, France: OpenReview.net. Available also from: <https://openreview.net/forum?id=SJU4ayYgl>.
- KUBARA, Kacper, 2020. *Machine Learning Tasks on Graphs* [Medium] [online]. 2020-09-28. [visited on 2020-11-03]. Available from: <https://towardsdatascience.com/machine-learning-tasks-on-graphs-7bc8f175119a>.
- MIKOLOV, Tomas; CHEN, Kai; CORRADO, Greg; DEAN, Jeffrey, 2013. Efficient Estimation of Word Representations in Vector Space [online] [visited on 2020-11-03]. Available from: <https://openreview.net/forum?id=idpCdOWtqXd60>.
- PEROZZI, Bryan; AL-RFOU, Rami; SKIENA, Steven, 2014. Deepwalk: Online learning of social representations. In: *Proceedings of the 20th ACM SIGKDD international conference on Knowledge discovery and data mining*, pp. 701–710.

VELIČKOVIĆ, Petar; DEAC, Andreea, 2020. *Opening Remarks*. Vienna. Available also from:
<https://slideslive.at/icml-2020/graph-representation-learning-and-beyond-grl>.
ICML 2020.

ZHOU, Jie; CUI, Ganqu; ZHANG, Zhengyan; YANG, Cheng; LIU, Zhiyuan; WANG, Lifeng;
LI, Changcheng; SUN, Maosong, 2019. Graph Neural Networks: A Review of Methods and Applications.
arXiv:1812.08434 [cs, stat] [online] [visited on 2020-11-03]. Available from arXiv: 1812.08434.